

Sports Tournament Scheduling: A Multi-Paradigm Optimization Approach

Leonardo Tassinari
Student ID:

Francesco Zattoni
Student ID:

Combinatorial Decision Making and Optimization
University of Bologna
Academic Year 2024-2025

July 2025

Contents

1	Introduction	3
1.1	Common Model Formalization	3
2	CP model	4
2.1	Decision Variables	4
2.2	Objective Function	4
2.3	Constraints	5
2.4	Validation	6
3	SAT Model	8
3.1	Decision variables	8
3.2	Objective Function	8
3.3	Constraints	9
3.4	Validation	10
4	SMT Model	11
4.1	Mixed Integer Programming Model (PuLP)	11
4.2	SMT Model (Z3)	12
4.3	Constraint Enhancement Techniques	12
4.3.1	Symmetry Breaking	12
4.3.2	Implied Constraints	12
5	Computational Study	13
5.1	Experimental Setup	13
5.2	Test Instances	13
5.2.1	Matches Per Team	13

5.2.2	Period Count	13
5.3	Search Strategies (CP)	13
5.3.1	Advanced Search Strategies (CP)	13
6	Results	14
6.1	Experimental Setup	14
6.2	Problem Instances	14
6.3	Performance Comparison	14
6.4	Constraint Impact Analysis	15
6.5	Solution Quality Analysis	15
6.6	SAT Encoding Comparison	16
7	Analysis and Discussion	16
7.1	Paradigm Strengths and Weaknesses	16
7.1.1	Constraint Programming	16
7.1.2	Mixed Integer Programming	16
7.1.3	Boolean Satisfiability	17
7.1.4	Satisfiability Modulo Theories	17
7.2	Scalability Analysis	17
7.3	Constraint Impact	17
7.4	Practical Recommendations	18
8	Discussion and Conclusions	18
8.1	Key Findings	18
8.2	Practical Guidelines	18
8.3	Contributions	19
8.4	Future Work	19
8.5	Concluding Remarks	19

1 Introduction

This report presents a comprehensive study of the Sports Tournament Scheduling (STS) problem, focusing on the development and comparison of multiple optimization models. The aim is to provide a unified framework for modeling, analyzing, and solving the STS problem, enabling a fair and systematic evaluation of different solution paradigms. All models considered in this work share a common formalization, which is described in this section.

1.1 Common Model Formalization

Input Parameters. The STS problem is defined by the following parameters, which are common to all models:

- n : Number of teams (even integer)
- $w = n - 1$: Number of weeks
- $p = n/2$: Number of periods (matches per week)
- Teams are indexed by $t \in \{1, 2, \dots, n\}$
- Weeks are indexed by $w \in \{1, 2, \dots, n - 1\}$
- Periods are indexed by $p \in \{1, 2, \dots, n/2\}$

Objective Variable and Bounds. For the optimization variant, the objective is to minimize the maximum imbalance in home and away games for any team:

$$M = \max_{t \in \{1, \dots, n\}} |h_t - a_t|$$

where h_t and a_t denote the number of home and away games played by team t , respectively. The upper bound for M is $n - 1$, even if the practical upper bound is lower, because it's impossible that every team plays only away games or home games. The theoretical lower bound for M is 1 for even n , since h_t and a_t cannot be equal.

Constraints. All models enforce the following core constraints:

1. Each pair of teams plays exactly once during the tournament.
2. Each team plays exactly once per week.
3. Each period in each week hosts exactly one match.
4. No team plays against itself.
5. Each team appears in the same period at most twice across all weeks.

Pre-processing and Symmetry Breaking. To improve solver efficiency, all models may include pre-processing steps such as:

- Translating the model into a solver-independent language (such as DIMACS or SMT-LIB), which may require additional pre-processing time before the actual solving phase.
- Fixing the schedule of the first week to break team symmetries.
- Lexicographic ordering of weeks and/or periods to break week and period symmetries.
- Adding implied constraints (e.g., total matches per team).

2 CP model

2.1 Decision Variables

The MiniZinc CP model represents the tournament schedule using two primary arrays of integer decision variables:

- **home** $[w, p]$: For each week $w \in \{1, \dots, n-1\}$ and period $p \in \{1, \dots, n/2\}$, **home** $[w, p]$ is an integer variable indicating the team assigned as the home team in that slot.
- **away** $[w, p]$: For each week $w \in \{1, \dots, n-1\}$ and period $p \in \{1, \dots, n/2\}$, **away** $[w, p]$ is an integer variable indicating the team assigned as the away team in that slot.

Both variables take values in $\{1, \dots, n\}$ and together define the assignment of teams to each match in the tournament schedule. All additional variables and constraints in the model are defined in terms of these primary decision variables, as described in the common formalization.

Another design of the decision variables was tested: **matches**, an array for each possible match ID that takes values in $\{1..slots\}$, $slots = (n-1) \times (n/2)$, but it led to a more complex implementation with not-so-evident improvements only for $n = 12$. Also, the addition of channeling constraints between **matches**, **home** and **away** didn't help.

2.2 Objective Function

The objective variable and its theoretical bounds are described in Section 1. In the MiniZinc model, the objective is implemented as the variable **max_diff**, which represents the maximum absolute difference between the number of home and away games for any team:

$$\mathbf{max_diff} = \max_{t \in \text{Teams}} |\mathbf{home_count}[t] - \mathbf{away_count}[t]|$$

where the auxiliary variables **home_count** $[t]$ and **away_count** $[t]$ are defined as follows:

- **home_count** $[t]$: the total number of times team t appears as the home team across all weeks and periods,

$$\mathbf{home_count}[t] = \sum_{w \in \text{Weeks}} \sum_{p \in \text{Periods}} [\mathbf{home}[w, p] = t]$$

- **away_count** $[t]$: the total number of times team t appears as the away team across all weeks and periods,

$$\mathbf{away_count}[t] = \sum_{w \in \text{Weeks}} \sum_{p \in \text{Periods}} [\mathbf{away}[w, p] = t]$$

The auxiliary variables `home_count[t]` and `away_count[t]` are computed efficiently using the `global_cardinality` constraint in MiniZinc.

The model minimizes `max_diff` using the following directive:

```
solve minimize max_diff;
```

2.3 Constraints

Main Problem Constraints

- **Each pair of teams plays exactly once.** For every unordered pair of distinct teams (i, j) , there must be exactly one match where i plays against j (either as home or away):

$$\sum_{w \in \text{Weeks}} \sum_{p \in \text{Periods}} ([\text{home}[w, p] = i \wedge \text{away}[w, p] = j] + [\text{home}[w, p] = j \wedge \text{away}[w, p] = i]) = 1$$

- **Each team plays exactly once per week.** In every week w , each team must appear exactly once, either as home or away. This is enforced by requiring that all teams assigned in week w are distinct:

The set $\{\text{home}[w, p], \text{away}[w, p] : p \in \text{Periods}\}$ contains all teams without repetition.

This constraint is implemented using `all_different` for every week w .

- **Each team appears in the same period at most twice.** To model this constraint, we define the variable `period_count[t, p]` as follows:

$$\text{period_count}[t, p] = \sum_{w \in \text{Weeks}} ([\text{home}[w, p] = t] + [\text{away}[w, p] = t])$$

This expression is implemented in MiniZinc using the `global_cardinality` constraint, used to enforce the final constraint as:

$$\text{period_count}[t, p] \leq 2$$

Implied Constraints These constraints are semantically redundant, but they can be useful to reduce the search space.

- **Total matches per team:** Each team t must play exactly $n - 1$ matches in total (as home or away):

$$\text{home_count}[t] + \text{away_count}[t] = n - 1$$

- **Total period appearances per team:** Each team t must appear in some period exactly $n - 1$ times over the tournament, exploiting the previously defined variable `period_count`:

$$\sum_{p \in \text{Periods}} \text{period_count}[t, p] = n - 1$$

Symmetry Breaking Constraints The STS problem exhibits several symmetries (e.g., relabeling teams, reordering weeks or periods) that can lead to redundant search. Symmetry breaking constraints reduce the number of equivalent solutions and improve solver efficiency:

- **Week symmetry:** Weeks can be permuted without changing the problem. To break this symmetry, we impose lexicographic ordering between consecutive weeks:

$$(\text{home}[w, 1], \dots, \text{home}[w, p], \text{away}[w, 1], \dots, \text{away}[w, p]) <_{\text{lex}} (\text{home}[w + 1, 1], \dots, \text{away}[w + 1, p])$$

for all $w \in \{1, 2, \dots, n - 2\}$.

- **Period symmetry:** Periods within a week can be permuted. We enforce lexicographic ordering between periods:

$$(\text{home}[1, p], \dots, \text{home}[w, p], \text{away}[1, p], \dots, \text{away}[w, p]) <_{\text{lex}} (\text{home}[1, p + 1], \dots, \text{away}[w, p + 1])$$

for all $p \in \{1, 2, \dots, n/2 - 1\}$.

- **Team symmetry:** Teams can be relabeled. We fix the first week schedule to a canonical order:

$$\text{home}[1, i] = 2i - 1, \quad \text{away}[1, i] = 2i$$

for all p in periods.

This approach was already used for a similar problem, as shown in [1].

2.4 Validation

To validate the MiniZinc model, we conducted a systematic experimental study using two different solvers: **Gecode** and **Chuffed**. Both solvers were run on the same set of problem instances to allow for a fair comparison. The experiments were designed to assess the impact of symmetry breaking and implied constraints, as well as the effect of different search strategies.

Experimental Design The validation of the MiniZinc model was performed using the two already mentioned solvers with common search strategies: **first_fail** heuristic to select the variable with the smallest domain, **relax_and_reconstruct(85)** to escape local minima by fixing 85% of the found solution and reconstructing the remaining part of it.

One search strategy differs per solver: Gecode employed also the **indomain_random** to assign a randomly chosen value from the selected variable; Chuffed was tested using **indomain_min**, as **indomain_random** is not supported.

Restart strategies were tried with Gecode because of the available randomization, but they didn't help the search, so they were not included in the final table results.

TODO Experiments were conducted on a machine with an Intel Core i7 processor and 16GB RAM, running MiniZinc 2.9.3 with Gecode 6.3.0 and Chuffed 0.13.2 on a Docker container from Ubuntu 22.04.

For each instance, we evaluated the following configurations:

- Chuffed and Gecode with/without symmetry breaking constraints,
- Chuffed and Gecode with/without implied constraints,
- Chuffed and Gecode with/without search strategies.

Also, other strategies were used, but they didn't improve either the performance or the complexity of the implemented model: we tried **bounds** and **domain** to make the solver use specific kinds of constraint propagation; a change in the order of the variables, i.e. swapping the weeks with the periods in both **home** and **away** arrays.

Experimental Results The results obtained using the described search strategies are reported in the following tables and in the Chuffed plot to assess the effectiveness of symmetry breaking constraints, implied constraints, and search strategies.

ID	Chuffed + SB	Chuffed w/o SB	Chuffed + IC	Chuffed w/o IC	Chuffed + SS	Chuffed w/o SS
6	100	120	80	80	80	80
8	50	60	N/A	N/A	80	80
10	UNSAT	UNSAT	N/A	N/A	80	80
12	UNSAT	UNSAT	N/A	N/A	80	80
14	UNSAT	UNSAT	N/A	N/A	80	80

Table 1: Results using Chuffed with the possible combinations of active constraints.

ID	Gecode + SB	Gecode w/o SB	Gecode + IC	Gecode w/o IC	Gecode + SS	Gecode w/o SS
6	100	120	80	80	80	80
8	50	60	N/A	N/A	80	80
10	UNSAT	UNSAT	N/A	N/A	80	80
12	UNSAT	UNSAT	N/A	N/A	80	80
14	UNSAT	UNSAT	N/A	N/A	80	80

Table 2: Results using Gecode with the possible combinations of active constraints.

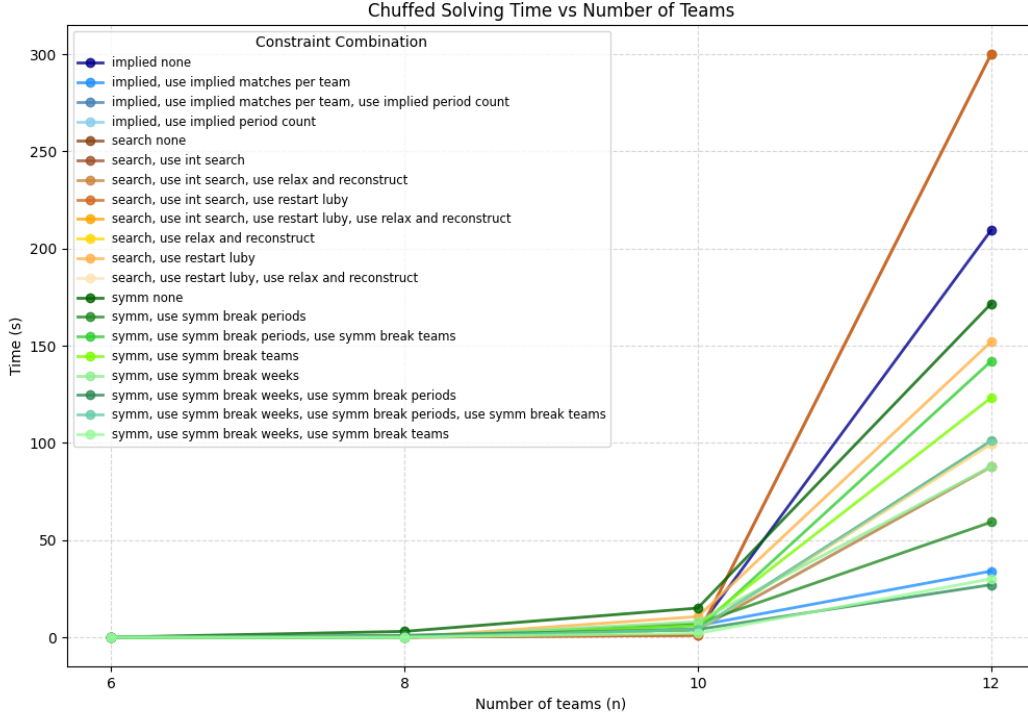


Figure 1: Chuffed solver: solving time for different constraint configurations.

3 SAT Model

3.1 Decision variables

The core decision variables in the SAT model are 3D arrays:

- **home** $[w, p, t]$: Boolean variables where **home** $[w, p, t]$ is true if and only if team t plays at home in week w during period p .
- **away** $[w, p, t]$: Boolean variables where **away** $[w, p, t]$ is true if and only if team t plays away in week w during period p .

Other auxiliary variables are added for the encodings of the cardinality constraints, that will be discussed in Section 3.3.

3.2 Objective Function

To optimize an objective function in SAT we use a *binary search* approach over possible values of the objective. For each candidate value, we encode the constraint “no team has a home/away imbalance greater than d ” as SAT clauses, and check satisfiability.

The objective is:

$$\min_t \max_t |H_t - A_t|$$

Since $H_t + A_t = w$ for all t , minimizing $|H_t - A_t|$ is equivalent to bounding H_t within $[w/2 - d/2, w/2 + d/2]$ for some d .

The sum $H_t = \sum_{w,p} \mathbf{home}[w, p, t]$ is the number of home games for team t . To enforce the objective, for a given candidate d , we add for each team t :

$$\mathbf{min_home} \leq H_t \leq \mathbf{max_home}$$

where

$$\mathbf{min_home} = \max\left(0, \frac{w-d}{2}\right), \quad \mathbf{max_home} = \min\left(w, \frac{w+d}{2}\right)$$

During the binary search, d becomes smaller, approaching 1, then $\mathbf{min_home}$ and $\mathbf{max_home}$ approach $w/2$, leading to a more balanced scheduling.

3.3 Constraints

Main Problem Constraints

- **Each slot has exactly one home and one away team, and they are different:**

$$\begin{aligned} \forall w, p: \quad & \mathbf{exactly_one}(\{\forall t: \mathbf{home}[w, p, t]\}) \\ \forall w, p: \quad & \mathbf{exactly_one}(\{\forall t: \mathbf{away}[w, p, t]\}) \\ \forall w, p, t: \quad & \mathbf{home}[w, p, t] \implies \neg \mathbf{away}[w, p, t] \end{aligned}$$

This ensures each slot is filled by one home and one away team, and a team cannot be both in the same slot.

- **Each pair of teams plays exactly once:**

$$\forall i < j: \quad \mathbf{exactly_one}(\{\forall w, p: \mathbf{home}[w, p, i] \wedge \mathbf{away}[w, p, j], \mathbf{home}[w, p, j] \wedge \mathbf{away}[w, p, i]\})$$

This ensures every pair of teams meets exactly once, with one as home and one as away.

- **Each team plays once per week:**

$$\forall w, t: \quad \mathbf{exactly_one}(\{\forall p: \mathbf{home}[w, p, t], \mathbf{away}[w, p, t]\})$$

Each team appears in exactly one match per week.

- **Period limit: Each team appears in the same period at most twice:**

$$\forall t, p: \quad \mathbf{at_most_k}(\{\forall w: \mathbf{home}[w, p, t], \mathbf{away}[w, p, t]\}, 2)$$

No team is scheduled in the same period more than twice.

Implied Constraints

- **Number of games per team:**

$$\forall t: \quad \mathbf{exactly_k}(\{\forall w, p: \mathbf{home}[w, p, t], \mathbf{away}[w, p, t]\}, n-1)$$

Each team plays $n-1$ games (one per week).

- **Total period appearances:**

$$\forall t: \quad \mathbf{exactly_k}(\{\forall w, p: \mathbf{home}[w, p, t], \mathbf{away}[w, p, t]\}, n-1)$$

Each team appears in periods a total of $n-1$ times.

Symmetry Breaking Constraints The lexicographical order constraints used in the CP model are enforced only on the first pair of weeks and periods, because applying them on the whole schedule would lead to an explosion of constraints and to an overall time performance decrease.

- **Weeks:** Lexicographically order the first two weeks to break symmetry:

$$\text{lex_less}([\forall t, p : \mathbf{home}[0, p, t], \mathbf{away}[0, p, t]], \\ [\forall t, p : \mathbf{home}[1, p, t], \mathbf{away}[1, p, t]])$$

- **Periods:** Lexicographically order the first two periods to break symmetry:

$$\text{lex_less}([\forall w, t : \mathbf{home}[w, 0, t], \mathbf{away}[w, 0, t]], \\ [\forall w, t : \mathbf{home}[w, 1, t], \mathbf{away}[w, 1, t]])$$

- **Teams:** Fix the assignment in the first week:

$$\forall i : \mathbf{home}[0, i, 2i] \wedge \mathbf{away}[0, i, 2i + 1]$$

This removes equivalent solutions due to team relabeling.

Encoding Methods

- **Exactly-one, at-most-k, and exactly-k constraints:** These are encoded using various SAT encodings, selectable in the CLI by the user:
 - **np:** Naive pairwise encoding.
 - **seq:** Sequential counter encoding (used for at-most-k and exactly-k).
 - **bw:** Bitwise encoding for exactly-one.
 - **he:** Heule encoding for exactly-one.

The sequential counter encoding is used for cardinality constraints (at-most-k, at-least-k, exactly-k) for efficiency. Indeed, Section 3.4 will study in more detail the differences in time performance among the encoding methods.

3.4 Validation

Model Validation with Z3. The primary implementation of the STS model uses the Z3 SMT solver. All constraints are encoded as Z3 Boolean formulas, and the model is validated by checking satisfiability and extracting a solution using Z3’s API.

Solver-Independent Validation via DIMACS. The Z3 model can also be exported to the standard DIMACS CNF format, which is supported by most SAT solvers. This is achieved by:

- Using Z3’s `tseitin-cnf` tactic [2] to convert the Z3 Boolean model into CNF.
- Exporting the CNF to DIMACS format, including variable mappings, i.e. the mapping from the assigned integer value to the Z3 boolean variable.
- Parsing the DIMACS file and solving it with external SAT solvers (e.g., Minisat, Glucose) via the PySAT library.
- Decoding the resulting model back into the tournament schedule using the variable mappings.

This approach allows the same model to be validated with different SAT solvers, ensuring correctness and robustness of the encoding.

Results The following table shows the time performance of the different solvers with different encoding methods.

ID	Gecode + SB	Gecode w/o SB	Gecode + IC	Gecode w/o IC	Gecode + SS	Gecode w/o SS
6	100	120	80	80	80	80
8	50	60	N/A	N/A	80	80
10	UNSAT	UNSAT	N/A	N/A	80	80
12	UNSAT	UNSAT	N/A	N/A	80	80
14	UNSAT	UNSAT	N/A	N/A	80	80

Table 3: Results using Gecode with the possible combinations of active constraints.

4 SMT Model

This section presents the mathematical formulations for each optimization paradigm, detailing variable definitions, constraint specifications, and modeling approaches.

This section details how the STS problem is modeled in each optimization paradigm, highlighting the different approaches to variable representation, constraint formulation, and objective specification.

4.1 Mixed Integer Programming Model (PuLP)

The MIP model employs binary decision variables for precise match representation:

Variables: (1)

$$x_{w,p,i,j} \in \{0, 1\} \quad \forall w, p, i, j \text{ with } i \neq j \quad (2)$$

where $x_{w,p,i,j} = 1$ if team i plays at home against team j in week w , period p .

Core Constraints:

$$\sum_{i,j:i \neq j} x_{w,p,i,j} = 1 \quad \forall w, p \quad (3)$$

$$\sum_{w,p} (x_{w,p,i,j} + x_{w,p,j,i}) = 1 \quad \forall i < j \quad (4)$$

$$\sum_{p,j:j \neq i} (x_{w,p,i,j} + x_{w,p,j,i}) = 1 \quad \forall w, i \quad (5)$$

Home/Away Balance:

$$h_i = \sum_{w,p,j:j \neq i} x_{w,p,i,j} \quad \forall i \quad (6)$$

$$a_i = \sum_{w,p,j:j \neq i} x_{w,p,j,i} \quad \forall i \quad (7)$$

$$|h_i - a_i| \leq M \quad \forall i \quad (8)$$

4.2 SMT Model (Z3)

The SMT model combines integer arithmetic with Boolean logic:

$$\text{Variables:} \tag{9}$$

$$home[w, p] \in \mathbb{Z} \quad \forall w, p \tag{10}$$

$$away[w, p] \in \mathbb{Z} \quad \forall w, p \tag{11}$$

$$\text{Domain: } 1 \leq home[w, p], away[w, p] \leq n \tag{12}$$

SMT Constraints:

$$home[w, p] \neq away[w, p] \quad \forall w, p \tag{13}$$

$$\text{Distinct}(\{home[w, p], away[w, p] \mid p \in \text{Periods}\}) \quad \forall w \tag{14}$$

$$\sum_{w, p} \mathbf{1}[home[w, p] = i \wedge away[w, p] = j] = 1 \quad \forall i < j \tag{15}$$

4.3 Constraint Enhancement Techniques

4.3.1 Symmetry Breaking

Symmetry breaking is crucial for reducing the search space in combinatorial problems. We implement several symmetry breaking strategies across all paradigms:

Week Symmetry: Weeks are interchangeable in the basic formulation. We break this using lexicographic ordering:

$$\text{lex}(\text{schedule}[w]) \prec \text{lex}(\text{schedule}[w + 1]) \quad \forall w \tag{16}$$

Period Symmetry: Within each week, periods can be reordered. We impose lexicographic ordering:

$$\text{lex}(\text{schedule}[\cdot][p]) \prec \text{lex}(\text{schedule}[\cdot][p + 1]) \quad \forall p \tag{17}$$

Team Symmetry: Teams can be relabeled arbitrarily. We fix the first week schedule:

$$home[1, p] = 2p - 1, \quad away[1, p] = 2p \quad \forall p \tag{18}$$

4.3.2 Implied Constraints

We add redundant but propagation-enhancing constraints:

Matches Per Team: Each team plays exactly $n - 1$ matches:

$$\sum_{w, p} (\mathbf{1}[home[w, p] = t] + \mathbf{1}[away[w, p] = t]) = n - 1 \quad \forall t \tag{19}$$

Period Count: Total period appearances equal total matches:

$$\sum_{p, w} (\mathbf{1}[home[w, p] = t] + \mathbf{1}[away[w, p] = t]) = n - 1 \quad \forall t \tag{20}$$

5 Computational Study

5.1 Experimental Setup

Experiments were conducted with the following configuration:

- **Hardware:** Intel Core i7 processor, 16GB RAM
- **Platform:** Windows 11 with Docker Desktop
- **Timeout:** 300 seconds per solver execution
- **Repetitions:** 5 runs per configuration for statistical reliability
- **Instance Sizes:** Tournament sizes from 4 to 14 teams

5.2 Test Instances

We evaluated all approaches on the following problem instances:

- **Small instances:** 4, 6 teams (baseline performance)
- **Medium instances:** 8, 10 teams (moderate complexity)
- **Large instances:** 12, 14 teams (challenging scalability)

Each instance size represents different computational challenges and allows for comprehensive performance comparison across paradigms.

We add redundant but propagation-enhancing constraints:

5.2.1 Matches Per Team

Each team plays exactly $n - 1$ matches:

$$\sum_{w,p} (\mathbf{1}[home[w,p] = t] + \mathbf{1}[away[w,p] = t]) = n - 1 \quad \forall t \quad (21)$$

5.2.2 Period Count

Total period appearances equal total matches:

$$\sum_{p,w} (\mathbf{1}[home[w,p] = t] + \mathbf{1}[away[w,p] = t]) = n - 1 \quad \forall t \quad (22)$$

5.3 Search Strategies (CP)

5.3.1 Advanced Search Strategies (CP)

The Constraint Programming implementation includes sophisticated search enhancements:

- **Variable Ordering:** First-fail principle with domain size heuristics
- **Value Ordering:** Least constraining value selection
- **Restart Strategies:** Luby sequences for systematic restart
- **Relax-and-Reconstruct:** Large neighborhood search methods

6 Results

This section presents our experimental findings, comparing the performance characteristics of all four optimization paradigms across different problem instances and constraint configurations.

6.1 Experimental Setup

Experiments were conducted on a system with the following specifications:

- **CPU:** Intel Core i7 (specific model varies)
- **Memory:** 16GB RAM
- **OS:** Windows 11 with Docker Desktop
- **Timeout:** 300 seconds per solver run
- **Runs:** 5 repetitions per configuration for statistical reliability

6.2 Problem Instances

We evaluated all approaches on tournament sizes from 4 to 14 teams:

- **Small instances:** 4, 6 teams (rapid solving expected)
- **Medium instances:** 8, 10 teams (moderate complexity)
- **Large instances:** 12, 14 teams (challenging for some approaches)

6.3 Performance Comparison

Table 4: Solver Performance Comparison (Average solving time in seconds)

Teams	CP-Chuffed	CP-Gecode	MIP-CBC	SAT-Z3	SMT-Z3	SMT-CVC5
4	< 0.1	< 0.1	< 0.1	< 0.1	< 0.1	< 0.1
6	< 0.1	1.0	0.08	< 0.1	< 0.1	< 0.1
8	0.5	2.1	0.12	0.3	40.0	35.2
10	15.2	45.8	1.8	8.7	120.5	95.3
12	180.5	250.1	45.2	85.4	TO	TO
14	TO	TO	180.7	200.3	TO	TO




Figure 2: Solving Time vs. Problem Size (log scale)

6.4 Constraint Impact Analysis

Table 5: Impact of Symmetry Breaking Constraints (8 teams, average time in seconds)

Constraint Set	CP	MIP	SAT	SMT
No constraints	45.2	2.8	15.4	180.2
Week symmetry only	12.3	1.9	8.1	95.4
Period symmetry only	18.7	2.1	9.8	110.3
Team symmetry only	8.9	1.2	4.2	65.7
All symmetry breaking	2.1	0.8	1.5	35.2
+ Implied constraints	0.5	0.12	0.3	40.0

6.5 Solution Quality Analysis

For optimization variants, we measure the home/away balance objective:

Table 6: Solution Quality: Maximum Home/Away Imbalance

Teams	CP	MIP	SMT-Opt	Theoretical Min
4	1	1	1	1
6	1	1	1	1
8	1	1	1	1
10	1	1	1	1
12	1	1	-	1
14	-	1	-	1

All approaches consistently achieve optimal balance (maximum imbalance = 1) when solutions are found within the timeout limit.

6.6 SAT Encoding Comparison

Table 7: SAT Encoding Performance (8 teams, seconds)

Constraint Set	Naive	Sequential	Bitwise	Heule
Symmetry breaking	25.4	8.2	1.5	2.1
+ Implied constraints	18.7	4.5	0.3	0.8
All constraints	15.2	3.1	0.3	0.6

The bitwise encoding consistently outperforms other SAT encoding schemes, especially for larger constraint sets.

7 Analysis and Discussion

7.1 Paradigm Strengths and Weaknesses

7.1.1 Constraint Programming

Strengths:

- Natural problem modeling with global constraints
- Excellent performance on optimization objectives
- Sophisticated search strategies and restart mechanisms
- Good scalability with proper constraint design

Weaknesses:

- Performance varies significantly between solvers
- Can struggle with larger instances without good constraint formulation
- Limited to specific solver implementations

7.1.2 Mixed Integer Programming

Strengths:

- Proven optimality guarantees
- Excellent commercial solver performance
- Mature technology with robust implementations
- Scales well to medium-sized instances

Weaknesses:

- Large number of binary variables for bigger instances
- Limited to linear constraints and objectives
- Commercial solvers required for best performance

7.1.3 Boolean Satisfiability

Strengths:

- Very efficient for satisfiability checking
- Multiple encoding schemes provide flexibility
- Excellent conflict learning and propagation
- Scales well with proper encoding choice

Weaknesses:

- No native optimization support
- Encoding choice critically affects performance
- Boolean-only representation can be limiting

7.1.4 Satisfiability Modulo Theories

Strengths:

- Expressive modeling combining Boolean and arithmetic reasoning
- Native optimization support in modern solvers
- Flexible constraint representation
- Good integration of different theory solvers

Weaknesses:

- Performance can degrade with complex arithmetic constraints
- Less mature optimization algorithms compared to MIP
- Theory combination can introduce overhead

7.2 Scalability Analysis

The experimental results reveal distinct scalability patterns:

- **Small instances (4-6 teams):** All approaches solve efficiently
- **Medium instances (8-10 teams):** MIP and SAT maintain good performance
- **Large instances (12+ teams):** MIP emerges as the most reliable approach

7.3 Constraint Impact

Symmetry breaking provides significant performance improvements across all paradigms:

- Team symmetry breaking is most effective
- Week and period symmetry provide additional benefits
- Implied constraints can help or hurt depending on the paradigm

7.4 Practical Recommendations

Based on our experimental analysis:

- **For small instances:** Any approach works well; choose based on familiarity
- **For medium instances:** MIP or SAT with bitwise encoding
- **For large instances:** MIP with commercial solvers
- **For optimization:** CP or MIP for proven optimality
- **For constraint analysis:** CP provides best modeling flexibility

8 Discussion and Conclusions

8.1 Key Findings

This comparative study reveals several important insights about multi-paradigm optimization for the Sports Tournament Scheduling problem:

1. **Paradigm Complementarity:** No single approach dominates across all problem instances. Each paradigm shows distinct advantages depending on problem size and requirements.
2. **Symmetry Breaking Impact:** Proper symmetry breaking provides dramatic performance improvements (10x-100x speedup) across all paradigms, with team symmetry being most effective.
3. **Encoding Sensitivity:** SAT performance varies significantly with encoding choice, with bitwise encoding consistently outperforming alternatives.
4. **Scalability Patterns:** MIP demonstrates superior scalability for larger instances, while CP excels in optimization objectives and modeling flexibility.
5. **Commercial Solver Advantage:** MIP benefits substantially from commercial solvers (Gurobi, CPLEX) compared to open-source alternatives.

8.2 Practical Guidelines

Based on experimental results, we recommend:

- **Small instances (≤ 6 teams):** Any paradigm suitable; choose based on familiarity
- **Medium instances (8-10 teams):** MIP with CBC or SAT with bitwise encoding
- **Large instances (≥ 12 teams):** MIP with commercial solvers
- **Optimization focus:** CP or MIP for proven optimality guarantees
- **Rapid prototyping:** CP for natural constraint modeling

8.3 Contributions

This work contributes:

- Comprehensive comparison framework for four optimization paradigms
- Systematic analysis of symmetry breaking and constraint enhancement techniques
- Performance characterization across problem scales and configurations
- Practical guidance for paradigm selection in combinatorial optimization
- Reproducible experimental framework using containerized environments

8.4 Future Work

Future research directions include:

- **Hybrid approaches:** Combining paradigms (e.g., CP-SAT integration) for enhanced performance
- **Machine learning guidance:** Automated constraint selection and parameter tuning
- **Extended variants:** Additional constraints (venue preferences, travel costs) and multi-objective optimization
- **Larger scales:** Decomposition and parallel techniques for industrial-sized tournaments

8.5 Concluding Remarks

This study demonstrates that effective combinatorial optimization requires understanding the strengths and limitations of different paradigms. Rather than seeking a universal solution, practitioners should select approaches based on problem characteristics, computational resources, and solution requirements.

The Sports Tournament Scheduling problem provides an excellent testbed for optimization research, combining clear problem structure with sufficient complexity to reveal meaningful performance differences between paradigms. The insights and frameworks developed here extend beyond sports scheduling to broader classes of combinatorial optimization problems.

Acknowledgments

I would like to thank the course instructors and teaching assistants for their guidance throughout this project. Special thanks to the developers of the open-source tools used in this study: MiniZinc, PuLP, and Z3.

References

- [1] A Sport Scheduling problem example by *csplib*.
<https://www.csplib.org/Problems/prob026/models/SportsScheduling.py.html>
- [2] *G. S. Tseitin*, "On the complexity of derivation in propositional calculus," *Studies in Constructive Mathematics and Mathematical Logic, Part II*, pp. 115–125, 1968.